



Simplifying the life-cycle management of HPC, data analytics and AI workflows

Jorge Ejarque (BSC), Rosa M Badia (BSC)

ACM Europe Summer School on HPC Computer Architectures
for AI and Dedicated Applications, Barcelona, July 5, 2023



This project has received funding from the European High-Performance Computing Joint Undertaking (JU) under grant agreement No 955558. The JU receives support from the European Union's Horizon 2020 research and innovation programme and Spain, Germany, France, Italy, Poland, Switzerland, Norway.
MCIN/AEI/10.13039/501100011033 and the European Union NextGenerationEU/PRTR (PCI2021-121957)

Agenda



14:00 – 15:25	Lectures: <ul style="list-style-type: none">- Overview of eFlows4HPC project- Integrating different computations in PyCOMPSs- HPC ready container images- Data pipelines and top-level workflow in TOSCA	Rosa M Badia
15:25 - 15:40	Bio Break	
15:40 - 16:55	Demo/ Hands-on session: How to build and deploy HPC Ready containers	Jorge Ejarque
16:55 - 17:00	Session conclusions	all presenters




OVERVIEW OF EFLOWS4HPC


HPC Environments



Current approach




Can we apply something like FaaS for Complex Workflows in HPC?




eFlows4HPC in a nutshell



- **Software tools stack that makes easier the development and management of complex workflows:**
 - Combine different frameworks
 - HPC, AI, data analytics
 - Reactive and dynamic workflows
 - Autonomous workflow steering
 - Full lifecycle management
 - Not just execution
 - Data logistics and Deployment
- **HPC Workflows as a Service:**
 - Mechanisms to make easier the use and reuse of HPC by wider communities
- **Architectural Optimizations:**
 - Selected HPC - AI Kernels Optimized for GPUs, FPGA, EPI
- **Validation Pillar's**
 - End-user workflows linked to CoEs



Pillar I: Manufacturing




6 - ACM Summer School, July 2023




Pillar I focuses on the construction of Digital Twins for the prototyping of complex manufactured objects:

- Integrating state-of-the-art adaptive solvers with machine learning and data-mining
- Contributing to the Industry 4.0 vision




5 July 2023

Pillar II: Climate



Data Analytics



Dynamic (AI-assisted) workflow




7 - ACM Summer School, July 2023

HPDA & ML/DL

- Perform climate predictions: temperature, precipitation or wind speed
- AI-assisted pruning of the ESM workflow
- Study of Tropical Cyclones (TC) in the North Pacific, with in-situ analytics



Pillar III: Urgent computing for natural hazards



Barcelona
Supercomputing
Center
Centro Nacional de Supercomputación

ETH zürich

8 - ACM Summer School, July 2023




Tsunami-HySEA GPU-based code



UNIVERSIDAD
DE MÁLAGA



NGI




5 July 2023

Pillar III explores the modelling of natural catastrophes:

- Earthquakes and their associated tsunamis shortly after such an event is recorded
- Use of AI to estimate intensity maps
- Use of DA and AI tools to enhance event diagnostics
- Areas: Mediterranean basin, Mexico, Iceland and Chile



HPCWaaS Overview



Project partners



Scuola
Internazionale
Superiore di
Studi Avanzati





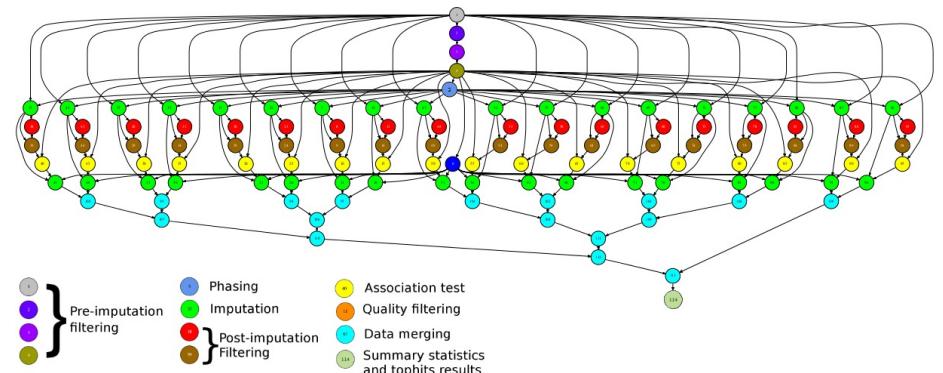
INTEGRATING DIFFERENT COMPUTATIONS IN PYCOMPSS

Main element: Workflows in PyCOMPSs



- Sequential programming, parallel execution
- General purpose programming language + annotations/hints
 - To identify tasks and directionality of data
- Builds a task graph at runtime that express potential concurrency
- Tasks can be sequential and parallel (threaded or MPI)
- Offers to applications the illusion of a shared memory in distributed system
 - The application can address larger data than storage space: support for Big Data apps
 - Support for persistent storage
- Agnostic of computing platform
 - Enabled by the runtime for clusters, clouds and container managed clusters

```
@task(c=INOUT)
def multiply(a, b, c):
    c += a*b
```



Barcelona
Supercomputing
Center

Centro Nacional de Supercomputación

PyCOMPSs syntax




- Use of decorators to annotate tasks and to indicate arguments directionality
- Small API for data synchronization

Tasks definition

```
@task(c=INOUT)
def multiply(a, b, c):
    c += a*b
```

```
initialize_variables()
startMulTime = time.time()
for i in range(MSIZE):
    for j in range(MSIZE):
        for k in range(MSIZE):
            multiply (A[i][k], B[k][j], C[i][j])
compss_barrier()
mulTime = time.time() - startMulTime
```



Synchronization



- Main program and tasks do not share the same memory spaces
- The synchronization `compss_wait_on` waits for tasks generating the parameter to be finished and moves the data from the remote node to the node where the main program is executed:

```
a = compute (b)
#compute is a task, here we can not check the value of a
...
a = compss_wait_on (a)
#here we can check the value of a
if a:
    ...
```

- Tasks can be also synchronized with a barrier

```
startMulTime = time.time()
for i in range(SIZE):
    compute (A[i], B[i])
compss_barrier()
multTime = time.time() - startMulTime
```

Other decorators: Tasks' constraints



- Constraints enable to define HW or SW features required to execute a task
 - Runtime performs the match-making between the task and the computing nodes
 - Support for multi-core tasks and for tasks with memory constraints
 - **Support for heterogeneity on the devices in the platform**

```
@constraint (MemorySize=6.0, ProcessorPerformance="5000", ComputingUnits="8")
@task (c=INOUT)
def myfunc(a, b, c):
    ...
```

```
@constraint (MemorySize=1.0, ProcessorType ="ARM", )
@task (c=INOUT)
def myfunc_other(a, b, c):
    ...
```

Failure Management



- Interface that enables the programmer to give hints about failure management

```
@task(file_path=FILE_INOUT, on_failure='CANCEL_SUCCESSORS',  
time_out='$task_timeout')  
def task(file_path):  
    ...  
    if cond :  
        raise Exception()
```


- Options: RETRY, CANCEL_SUCCESSORS, FAIL, IGNORE
- Implications on file management:
 - I.e., on IGNORE, output files: are generated empty
- **Possibility of ignoring part of the execution of the workflow, for example if a task fails in an unstable device**
- **Opens the possibility of dynamic workflow behaviour depending on the actual outcome of the tasks**

- Tasks can raise exceptions

```
@task(file_path=FILE_INOUT)  
def comp_task(file_path):  
    ...  
    raise COMPSSException("Exception raised")
```

- Combined with groups of tasks enables to cancel the group of tasks on the occurrence of an exception

```
def test_cancellation(file_name):  
    try:  
        with TaskGroup('failedGroup'):  
            long_task(file_name)  
            long_task(file_name)  
            executed_task(file_name)  
            comp_task(file_name)  
            cancelledTask(FILE_NAME);  
            cancelledTask(FILE_NAME)  
  
    except COMPSSException:  
        print("COMPSSException caught")  
        write_two(file_name)
```



Other decorators: linking with other programming models



- A task can be more than a sequential function
 - A task in PyCOMPSs can be sequential, multicore or multi-node
 - External binary invocation: wrapper function generated automatically
 - Supports for alternative programming models: MPI and OmpSS
- Additional decorators:
 - `@binary(binary="app.bin")`
 - `@mpi(binary="mpiApp.bin", runner="mpirun", processes=8)`
 - `@ompss(binary="ompssApp.bin")`
- Can be combined with the `@constraint` and `@implement` decorators


```
@binary(binary="app.bin", workingDir="/myApp")
@task()
def func(l):
    pass
```

Support for MPI tasks



- Resource manager aware of multi-node tasks

```
@mpi (runner="mpirun", processes= "32", processes_per_node=8)
@task (returns=int, stdOutFile=FILE_OUT_STDOUT, stdErrFile=FILE_OUT_STDERR)
def nems(stdOutFile, stdErrFile):
    pass
```



MPMD applications



- The `@mpmd_mpi` decorator can be used to define Multiple Program Multiple Data (MPMD) MPI tasks

```
@mpmd_mpi(runner="mpirun", working_dir = {{working_dir_exe}},
            programs=[{binary="fesom.x", processes = "$FESOM_PROCS" },
                      {binary="oifs", args="-v ecmwf -e awi3", processes = "$OIFS_PROCS" },
                      {binary="rnfma", processes = "$RNFMA_PROCS"}])
@task(log_file={Type:FILE_OUT, StdOutputStream:STDOUT}, working_dir_exe= DIRECTORY_INOUT)
def esm_simulation(log_file, working_dir_exe):
    pass
```

- As a result of the `@mpmd_mpi` annotation, the following commands will be generated:

```
> cd working_dir_exe; mpirun -n $FESOM_PROCS fesom.x : \
    -n $OIFS_PROCS oifs -v ecmwf -e awi3 : -n $RNFMA_PROCS rnfma
```

Tasks in container images



- Goal: enable tasks embedded in container images
 - @container decorator can be used together with the task annotation
 - Also support for user-defined tasks


```
@container(engine="DOCKER", image="ubuntu")
@binary(binary="ls")
@task()
def task_binary_empty():
    pass
```

```
@container(engine="DOCKER", image="compss/compss")
@task(returns=1, num=IN, in_str=IN, fin=FILE_IN)
def task_python_return_str(num, in_str, fin):
    print("Hello from Task Python RETURN")
    print("- Arg 1: num -- " + str(num))
    print("- Arg 1: str -- " + str(in_str))
    print("- Arg 1: fin -- " + str(fin))
    return "Hello"
```


PyCOMPSs runtime



- Runtime deployed as a distributed master-worker
- All data scheduling decisions and data transfers are performed by the runtime
- Support for elasticity



Interfaces to integrate HPC/DA/ML



- **Goal:**
 - Reduce the required glue code to invoke multiple complex software steps
 - Developer can focus in the functionality, not in the integration
 - Enables reusability
- **Two paradigms:**
 - Software invocation
 - Data transformations

```
#workflow steps defined as tasks
@data_transformation (input_data, transformation_description)
@software (invocation_description)
def data_analytics (input_data, result):
    pass

#workflow body
simulation (input_cfg, sim_out)
data_analytics (sim_out, analysis_result)
ml_training (analysis_result, ml_model)
```



Software Invocation description



Admin/user code

```
{  
  "type": "mpi",  
  "properties": {  
    "runner": "mpirun",  
    "processes": "$SW_PROCS",  
    "binary": "mpi_software.x",  
    "params": "-d {{param}}",  
    "working_dir": "{{{working_dir}}}",  
    "prolog": {  
      "binary": "mkdir",  
      "params": "{{{working_dir}}}"},  
    "epilog": {  
      "binary": "tar",  
      "params": "zcvf {{out_tgz}} {{{working_dir}}}",  
    },  
    "constraints": {  
      "computing_units": $SW_THREADS  
    }  
}
```

Software invocation
description
Stored in software catalog



- Converts a Python function of a software invocation to a PyCOMPSs task
- Takes information from the description in json
- Enables reuse in multiple workflows

Data transformations




- A data transformation changes the data without requiring extra programming from the developer

Admin/user code


```
def load_blocks_rechunk(blocks, shape, block_size, new_block_size):  
    ...  
    SnapshotMatrix = load_blocks_array (final_blocks, shape, block_size);  
    return SnapshotMatrix
```

```
@dt("blocks") load_blocks_rechunk, shape=expected_shape, block_size=simulation_block_size,  
    new_block_size=desired_block_size, is_workflow=True)  
@software(config_file = SW_CATALOG + "/dislib/dislib.json")  
def rSVD(blocks, desired_rank=30):  
    u,s = rsda(blocks, desired_rank, A_row_chunk_size, A_column_chunk_size)  
    return u
```

User code



Pillar I: Integration of HPC and data analysis workflow




```

fom.json
{
    "execution" : {
        "type": "task"
    },
    "constraints" : {
        "computing_units": "$KRATOS_CUS"
    },
    "parameters" : {
        "returns" : 1,
        "model" : "IN",
        "parameters" : "IN",
        "sample" : "IN"
    }
}


@software(config_file = SW_CATALOG+/kratos/fom.json")
def execute_FOM_instance(model, parameters, sample):
    import KratosMultiphysics
    from kratos_simulations import RunROM_SavingData
    load_ROM(rom)
    current_model = KratosMultiphysics.Model()
    model.Load("ModelSerialization",current_model)
    del(model)
    current_parameters = KratosMultiphysics.Parameters()
    parameters.Load("ParametersSerialization",current_parameters)
    del(parameters)
    simulation = RunROM_SavingData(current_model,current_parameters,sample)
    simulation.Run()
    return simulation.GetSnapshotsMatrix()

```

Pillar I: Integration of HPC and data analytics in a single workflow



Pillar I: Integration of HPC and data analytics in a single workflow




```
dislib.json
{
    "execution": {
        "type": "workflow"
    }
}
```

Method that invokes other
PyCOMPSs tasks inside

```
@dt("blocks", load_blocks_rechunk, shape=expected_shape,
block_size=simulation_block_size, new_block_size=desired_block_size,
is_workflow=True)
@software(config_file = SW_CATALOG + "/py-dislib/dislib.json")
def rSVD(blocks, desired_rank=30):
    from dislib_randomized_svd import rsvd
    u,s = rsvd(blocks, desired_rank, A_row_chunk_size, A_column_chunk_size)
    return
```

Pillar I: Integration of HPC and data analytics in a single workflow



```
import dislib as ds
```

```
@dt("blocks", load_blocks_rechunk, shape=expected_shape,
block_size=simulation_block_size,new_block_size=desired_
is_workflow=True)
@software(config_file = SW_CATALOG + "/dislib/dislib.json")
def rsrd(blocks, desired_rank):
    k = desired_rank
    ...
    Y = A @ Omega
    Q,R = my_qr(Y._blocks)
    Q=load_blocks_rechunk([Q], ...)
    ...
    return Q, R
```


```
@software(config_file=kratos.json)
def ExecuteInstance (model, parameters, Cases, instance):
    ...
    current_parameters = ....
    ...
    simulation = Get('simulation')
    simulation.Run()
    return simulation
```

```
for instance in range (0,TotalNumberOfCases):
    blocks.append(ExecuteInstance(model, parameters, pars, instance))
    ...
    U, s = rSVD(A, desired_rank)
    ...
    ...
```



HPC READY CONTAINER IMAGES

FaaS vs HPCWaaS




Similarities

- Easy to use for final user
- **Automate deployment & execution**
- Data integration
- **Containers**

Differences


- Restrictions
- Deployment and Execution Complexity
- **Performance**



Containers and HPC




Standard container image creation




- **Simplicity for deployment**
 - Just pull or download the image
- **Trade-Off performance/portability**
 - Architecture Optimizations
- **Accessing Hardware from Containers**
 - MPI Fabric /GPUs
- **Host-Container Version Compatibility**

HPC Ready Containers



- **Methodology to allow the creation containers for specific HPC system**
 - Leverage HPC and Multi-platform container builders
- **It is tight to do by hand but let's automate!**

Container Image Creation Service



Container Image Creation Service



- Web Interface

The screenshot shows the 'New Container Image Build Request' page. At the top, there's a navigation bar with 'eFlows4HPC' logo, 'Dashboard', and 'Logout'. On the left, a sidebar has 'Home', 'Builds', 'Images', and 'Account' options. The main area has 'Machine Description' fields for 'System Platform' (dropdown), 'Processor Architecture' (text input), 'Container Engine' (dropdown), 'MPI version (Optional)' (text input), and 'GPU runtime (Optional)' (text input). Below that is 'Workflow Reference' with 'Workflow Name' and 'Sub-workflow Name' fields. At the bottom is a large blue 'Build' button.

- REST Interface and CLI

The screenshot shows a POST request to the '/build/' endpoint. The request body is a JSON object:

```
{  
  "machine": {  
    "platform": "linux/amd64",  
    "architecture": "rome",  
    "container_engine": "singularity"},  
  "workflow": "minimal_workflow",  
  "step_id": "wordcount",  
  "force": False  
}
```

The response is:


HTTP/1.1 200 OK
Content-Type: application/json

```
{  
  "id": "<creation_id>"  
}
```

At the bottom, there's a terminal window showing the command to run the CLI:

```
localhost:~/image_creation> ./cic_cli <user> <token> https://<image_creation_url> build <request.json>  
Response:  
{ "id": "f1f4699b-9048-4ecc-aff3-1c689b855adc" }
```

Example: Pillar I for MN4




Spack Environment Concretization



```
"machine": {  
    "platform": "linux/amd64",  
    "architecture": "skylake_avx512",  
    "container_engine": "singularity",  
    "mpi": intel-mpi@2018.4  
},
```

```
spack:  
  specs:  
    - compss  
    - py-dislib  
    - kratos@9.1.4 apps=LinearSolversApplication, \  
      FluidDynamicsApplication, \  
      StructuralMechanicsApplication, \  
      ConvectionDiffusionApplication
```




```
spack:  
  specs:  
    - compss  
    - py-dislib  
    - kratos@9.1.4 apps=LinearSolversApplication, \  
      FluidDynamicsApplication, \  
      StructuralMechanicsApplication, \  
      ConvectionDiffusionApplication  
    - intel-mpi@2018.4  
concretization: together  
view: /opt/view  
packages:  
  all:  
    target: ['skylake_avx512']
```


Performance




Kratos Multiphysics (shared memory)



FESOM2 (MPI)



Tsunami-HySEA CTE-Power



Take away message



- **HPC ready container images**
 - Advantages:
 - Reduce the deployment complexity
 - Performance close to bare metal
 - Drawbacks:
 - Larger building times compared to OS packages
 - Require more compilations
 - Configure spack binary caches
 - Less portable
 - CPU architecture, mpi, gpu versions



DATA PIPELINES AND TOP-LEVEL WORKFLOW IN TOSCA

Data pipelines: motivation



eFlows4HPC aims to deliver a workflow software stack and an additional set of services to enable the integration of HPC simulations and modelling with big data analytics and machine learning in scientific and industrial applications. The software stack will allow for the creation of innovative adaptive workflows that efficiently use computing resources considering novel storage solutions.

- Computations require (lots) of data
- Data Logistics Service (DLS): fuel the scientific calculations with required data
- DLS pipelines describe how the data are moved

Data pipelines motivation



- Assumptions about data:
 - Right data
 - Right place
 - Right time
 - → Often proven wrong! (hacks! “legacy solutions”)
- Requirement to publish the data
 - → Often forgotten
- Data Pipelines (as part of scientific workflows):
 - Formalization
 - Reproducibility (FAIR)
 - Portability of the workflows
- In eFlows4HPC, data pipelines implemented with Apache airflow

Data Catalogue and Data Logistics Service



Data Catalogue:

- Lists datasets used and created by the workflow according to FAIR principles
- Provides metadata to make data movement pipelines more generic

Data Pipelines:



- Formalization of data movements for transparency and reusability
- Stage-in/out, image transfer

Data Logistics Services (DLS):

- Performs the execution of data pipelines at deployment and execution time

Production Ready Services:


- <https://datacatalogue.eflows4hpc.eu>
- <https://datalogistics.eflows4hpc.eu/>



Top-level workflows approach




- Requires a description for workflow lifecycle management
 - TOSCA:
 - Model to describe cloud application topologies and its lifecycle orchestration
- Interface for deploying and running the workflows
 - HPCWaaS:
 - Development and Deployment (Alien4Cloud)
 - Execution (HPCWaaS API)




TOSCA Modelization



Topology of the different components involved in the Workflow lifecycle management



Application deployment (done once)



Workflow publication and users authorize



The screenshot shows the eFlows4HPC interface for managing a workflow named 'pillar_I'. The top navigation bar includes icons for Applications and Catalog, and the current page title 'pillar_I'. A green cartoon character is positioned in the top right corner.

pillar_I

Description: pillar_I

ID: pillar_I

Creation date: Thu, May 4, 2023 12:05 PM

Update date: Thu, May 4, 2023 12:05 PM


Actions:

- Versions
- Environments
- Variables
- Users and Groups
- Delete

Tags:

hpcwaas-workflows	exec_job	
hpcwaas-authorized-users	jorge, loic,jedrzej	

End-user workflow execution (multiple executions)



Links



- Project website: eflows4hpc.eu
 - From software tab, access to github and documentation
- Project github: github.com/eflows4hpc
 - Specific repo for the ISC-HPC tutorial



eFlows4HPC

Enabling dynamic and Intelligent workflows
in the future EuroHPC ecosystem

www.eFlows4HPC.eu



@eFlows4HPC



eFlows4HPC Project



This project has received funding from the European High-Performance Computing Joint Undertaking (JU) under grant agreement No 955558. The JU receives support from the European Union's Horizon 2020 research and innovation programme and Spain, Germany, France, Italy, Poland, Switzerland, Norway.